

LOAN-APPROVAL PREDICTION PROJECT IN MACHINE LEARNING

Parth Mishra

OBJECTIVE

The Loan Approval Prediction project aims to automate the loan approval process using machine learning algorithms. It analyzes various applicant factors such as income, credit score, loan amount, employment type, and financial history to predict whether a loan should be approved or rejected. The model helps banks and financial institutions make faster, unbiased, and data-driven decisions, reducing human errors.

By leveraging historical data, it enhances accuracy, efficiency, and fairness in the approval process. The ultimate goal is to streamline operations, minimize manual workload, and improve customer satisfaction through reliable predictions

KEY INSIGHTS

- 1. Applicants with higher income levels and strong credit histories are more likely to get their loans approved.**
- 2. A high loan amount relative to income or poor credit score reduces approval chances.**
- 3. Employment stability and lower debt-to-income ratios play a major role in loan eligibility.**
- 4. The model achieves high accuracy in predicting approvals, minimizing manual decision errors.**
- 5. Using ML helps speed up processing while ensuring fair and data-driven loan evaluations**

SUMMARY

This Loan Approval Prediction project uses machine learning to automate the loan approval process. It analyzes applicant data such as income, credit score, loan amount, and employment type to predict loan eligibility. The goal is to make loan decisions faster, more accurate, and unbiased, reducing manual effort and human error. By using predictive modeling, the system helps financial institutions improve efficiency and fairness in loan processing.

RECOMMENDATIONS

- 1. Integrate real-time data such as updated credit scores and income verification to improve prediction accuracy.**
- 2. Balance the dataset by handling class imbalance to avoid bias toward loan rejections or approvals.**
- 3. Use feature importance analysis to focus on the most impactful factors like income, credit score, and loan amount.**
- 4. Regularly retrain the model with new data to adapt to changing financial and economic trends.**
- 5. Implement a dashboard or alert system for loan officers to interpret predictions transparently.**
- 6. Apply explainable AI (XAI) tools to make the decision making process more understandable and trustworthy.**
- 7. Continuously monitor performance metrics (accuracy, precision, recall) to maintain model reliability**

```

import pandas as pd

df = pd.read_csv(r"C:\Users\parth\OneDrive\Desktop\loan_approval_dataset.csv")
df.head()

      loan_id  no_of_dependents  education  self_employed
income_anum \
0           1                  2    Graduate        No
9600000
1           2                  0  Not Graduate       Yes
4100000
2           3                  3    Graduate        No
9100000
3           4                  3    Graduate        No
8200000
4           5                  5  Not Graduate       Yes
9800000

      loan_amount  loan_term  cibil_score
residential_assets_value \
0      29900000          12            778
2400000
1      12200000          8             417
2700000
2      29700000         20            506
7100000
3      30700000          8             467
18200000
4      24200000         20            382
12400000

      commercial_assets_value  luxury_assets_value
bank_asset_value \
0                  17600000            22700000
8000000
1                  22000000            8800000
3300000
2                  45000000            33300000
12800000
3                  33000000            23300000
7900000
4                  82000000            29400000
5000000

      loan_status
0      Approved
1     Rejected
2     Rejected
3     Rejected
4     Rejected

```

```

import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score,
    f1_score, confusion_matrix, classification_report,
    roc_auc_score, RocCurveDisplay
)
RND = 42

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")

from sklearn.model_selection import train_test_split
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.pipeline import Pipeline

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier,
GradientBoostingClassifier

from sklearn.metrics import (
    accuracy_score, precision_score, recall_score, f1_score,
    roc_auc_score, classification_report, confusion_matrix,
    roc_curve, auc
)
RND = 42

# 2. Clean column names & quick info

# make names: lowercase, strip spaces, replace inner spaces with _
df.columns = [str(c).strip().lower().replace(" ", "_") for c in

```

```

df.columns]

print("Columns:", df.columns.tolist())
print("\nData info:")
display(df.info())

Columns: ['loan_id', 'no_of_dependents', 'education', 'self_employed',
'income_annum', 'loan_amount', 'loan_term', 'cibil_score',
'residential_assets_value', 'commercial_assets_value',
'luxury_assets_value', 'bank_asset_value', 'loan_status']

Data info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4269 entries, 0 to 4268
Data columns (total 13 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   loan_id          4269 non-null    int64  
 1   no_of_dependents 4269 non-null    int64  
 2   education        4269 non-null    object  
 3   self_employed    4269 non-null    object  
 4   income_annum    4269 non-null    int64  
 5   loan_amount      4269 non-null    int64  
 6   loan_term        4269 non-null    int64  
 7   cibil_score      4269 non-null    int64  
 8   residential_assets_value 4269 non-null    int64  
 9   commercial_assets_value 4269 non-null    int64  
 10  luxury_assets_value 4269 non-null    int64  
 11  bank_asset_value 4269 non-null    int64  
 12  loan_status      4269 non-null    object  
dtypes: int64(10), object(3)
memory usage: 433.7+ KB

```

None

```

#Identify columns
# our target is 'loan_status', but we make it robust:
candidates = [c for c in df.columns if c in (
    "loan_status", "status", "approved", "approval",
    "loan_decision", "loan_approved"
)]
if not candidates:
    candidates = [c for c in df.columns if "loan" in c and "status" in c]
if candidates:
    target_col = candidates[0]
else:
    target_col = df.columns[-1]  # fallback: last column

```

```

print(f"Selected target column: {target_col!r}")
print("Unique target values (sample):", df[target_col].unique()[:10])

# --- map Approved / Rejected etc. to binary ---
y_raw = df[target_col].astype(str).str.strip().str.lower()

def map_to_binary(s):
    if pd.isna(s):
        return np.nan
    if s in {"y", "yes", "approved", "accept", "accepted", "1", "true", "t"} or "approve" in s:
        return 1
    if s in {"n", "no", "rejected", "reject", "0", "false", "f", "deny", "denied"} or "reject" in s:
        return 0
    # if numeric 0/1
    try:
        v = float(s)
        if v == 1.0: return 1
        if v == 0.0: return 0
    except Exception:
        pass
    # default guess
    if "yes" in s or "approve" in s:
        return 1
    return 0

y = y_raw.apply(map_to_binary)

print("\nTarget distribution after mapping:")
display(y.value_counts(dropna=False))

# drop rows with missing target (should be none)
mask_na = y.isna()
if mask_na.sum() > 0:
    print(f"Dropping {mask_na.sum()} rows with missing target")
    df = df.loc[~mask_na].reset_index(drop=True)
    y = y.loc[~mask_na].reset_index(drop=True)
else:
    y = y.reset_index(drop=True)

Selected target column: 'loan_status'
Unique target values (sample): ['Approved' 'Rejected']

Target distribution after mapping:

loan_status
1    2656
0    1613
Name: count, dtype: int64

```

```

# Quick EDA & Visualizations
from IPython.display import display

# Missing values
missing = df.isnull().sum().sort_values(ascending=False)
print("Columns with missing values:")
display(missing[missing > 0])

if (missing > 0).any():
    plt.figure(figsize=(10,4))
    missing[missing > 0].plot(kind="bar")
    plt.title("Missing values per column")
    plt.ylabel("Count")
    plt.tight_layout()
    plt.show()
else:
    print("No missing values found. Skipping missing-value plot.")

# Class balance
plt.figure(figsize=(5,4))
sns.countplot(x=pd.Series(y))
plt.title("Target distribution (0=Rejected, 1=Approved)")
plt.xlabel("Loan status")
plt.ylabel("Count")
plt.tight_layout()
plt.show()

# numeric / categorical columns (excluding id)
id_like = [c for c in df.columns if c.endswith("_id") or c ==
"loan_id" or c.startswith("id")]
print("ID-like columns (will be dropped later):", id_like)

numeric_cols = df.select_dtypes(include=[np.number]).columns.tolist()
numeric_cols = [c for c in numeric_cols if c not in id_like]
categorical_cols = df.select_dtypes(include=["object", "category",
"bool"]).columns.tolist()
categorical_cols = [c for c in categorical_cols if c != target_col]

print("Numeric cols:", numeric_cols)
print("Categorical cols:", categorical_cols)

# Histograms for numeric
if numeric_cols:
    df[numeric_cols].hist(bins=25, figsize=(14,6))
    plt.suptitle("Numeric feature distributions")
    plt.tight_layout(rect=[0,0,1,0.95])
    plt.show()

# Bar charts for categoricals
for c in categorical_cols:

```

```

plt.figure(figsize=(6,3))
df[c].value_counts().plot(kind="bar")
plt.title(f"Top categories: {c}")
plt.tight_layout()
plt.show()

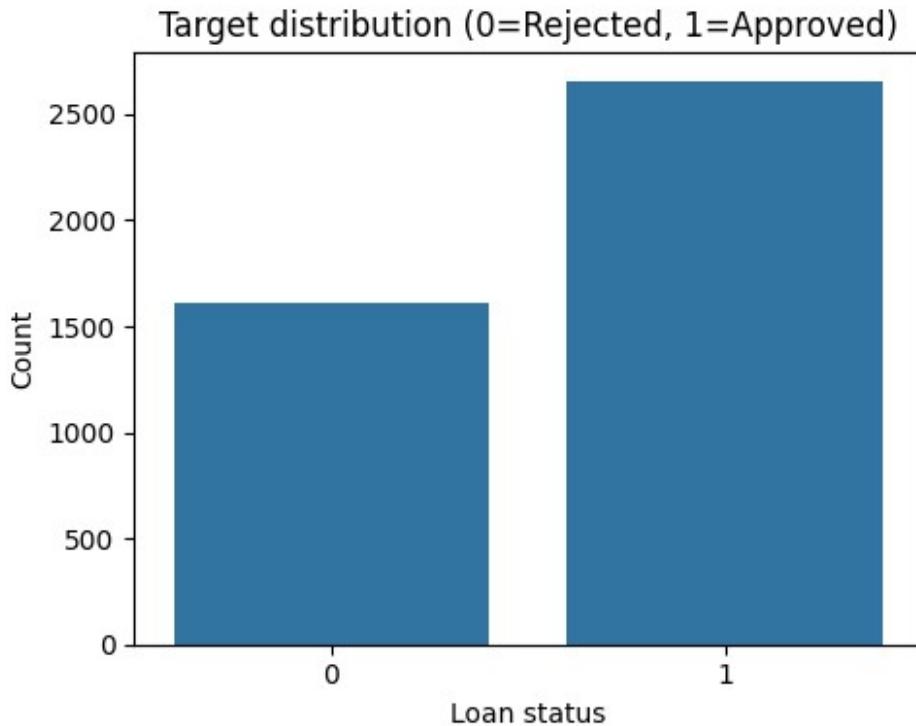
# Correlation heatmap (numeric + target)
corr_df = df[numeric_cols].copy()
corr_df["__target__"] = y.values.astype(float)

plt.figure(figsize=(10,8))
sns.heatmap(corr_df.corr(), annot=False, cmap="coolwarm", center=0)
plt.title("Correlation matrix (numeric features + target)")
plt.tight_layout()
plt.show()

Columns with missing values:
Series([], dtype: int64)

□ No missing values found. Skipping missing-value plot.

```

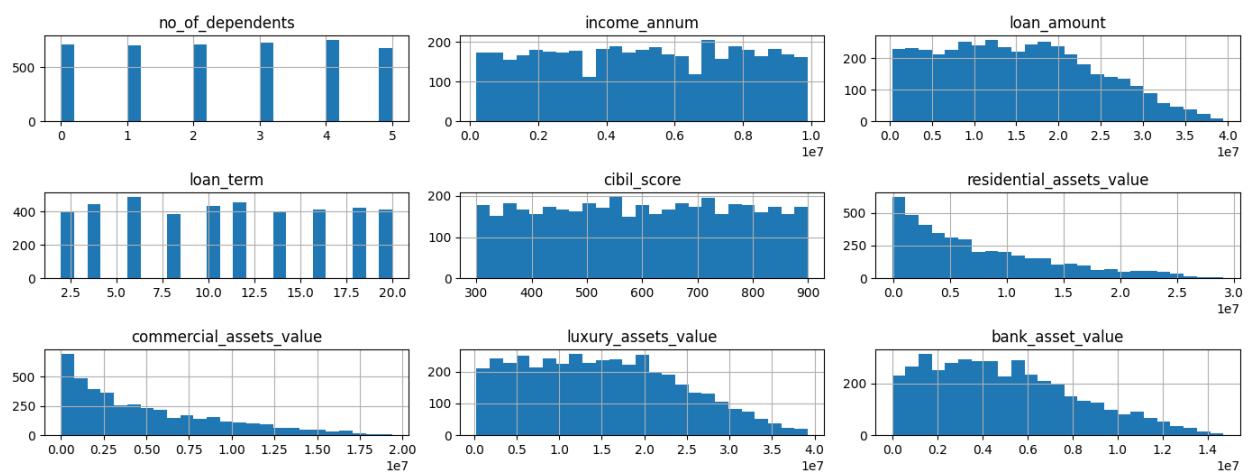


```

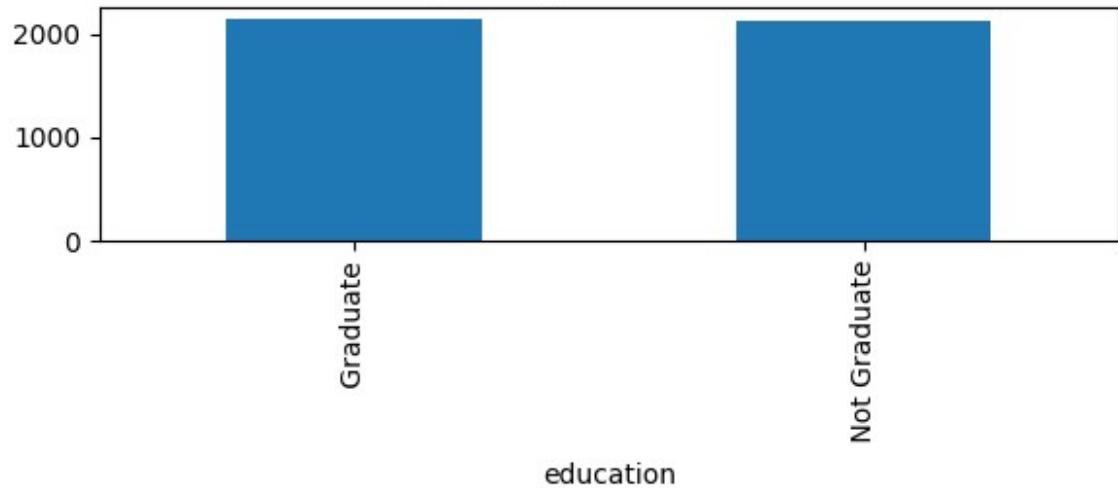
ID-like columns (will be dropped later): ['loan_id']
Numeric cols: ['no_of_dependents', 'income_annum', 'loan_amount',
'loan_term', 'cibil_score', 'residential_assets_value',
'commercial_assets_value', 'luxury_assets_value', 'bank_asset_value']
Categorical cols: ['education', 'self_employed']

```

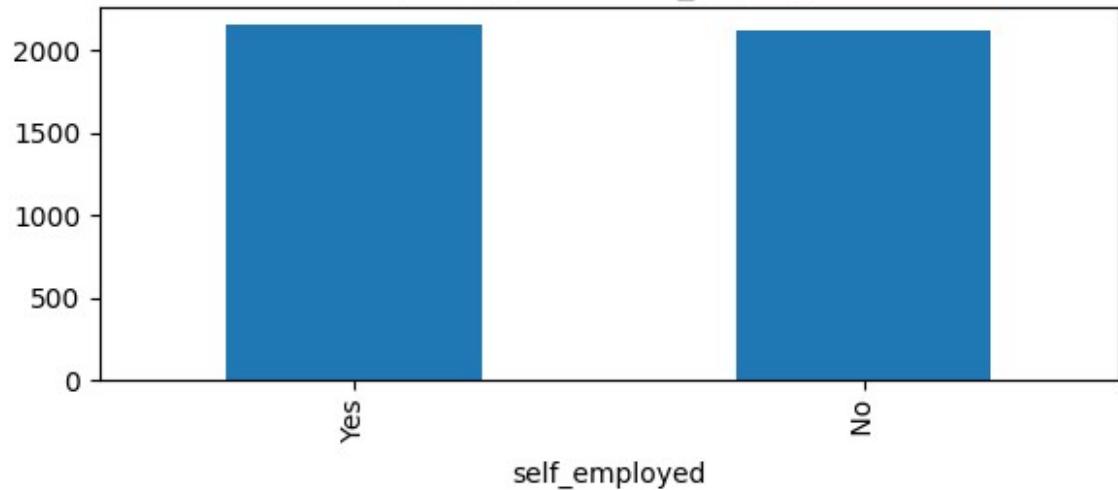
Numeric feature distributions

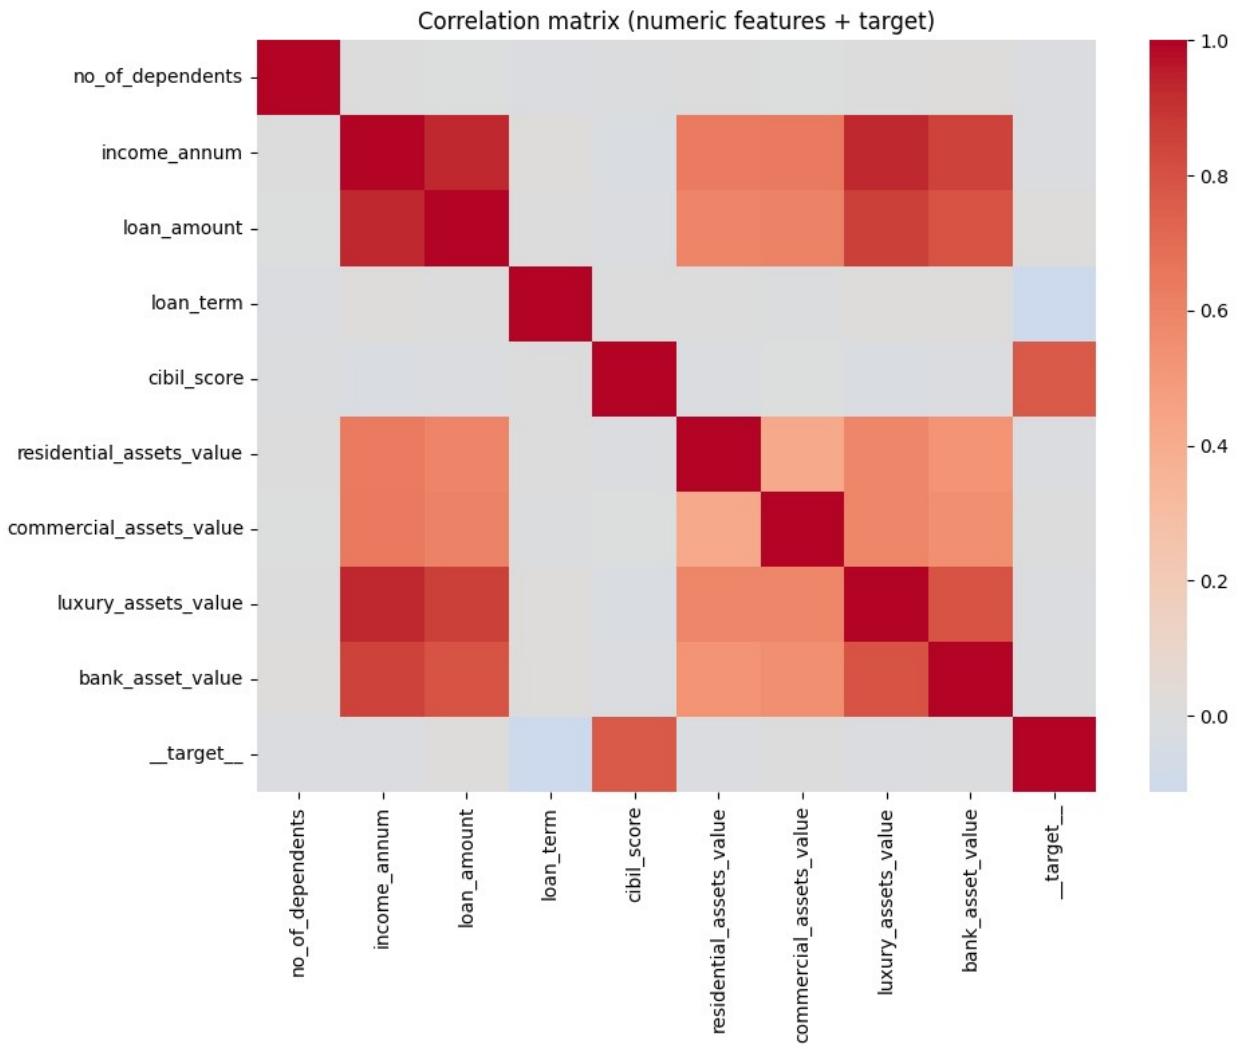


Top categories: education



Top categories: self_employed





```
# Train-Test Split
X = df.drop(columns=[target_col], errors="ignore")
y = y.astype(int)

# drop id columns
for c in id_like:
    if c in X.columns:
        X = X.drop(columns=c)

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=RND, stratify=y
)

print("Train shape:", X_train.shape)
print("Test shape :", X_test.shape)

Train shape: (3415, 11)
Test shape : (854, 11)
```

```

# Pre-processing (ColumnTransformer)
num_cols = X.select_dtypes(include=np.number).columns
cat_cols = X.select_dtypes(include=["object", "category",
"bool"]).columns

print("Numeric columns used:", list(num_cols))
print("Categorical columns used:", list(cat_cols))

numeric_pipe = Pipeline([
    ("imputer", SimpleImputer(strategy="median")),
    ("scaler", StandardScaler())
])

# NOTE:
# then change sparse_output=False
categorical_pipe = Pipeline([
    ("imputer", SimpleImputer(strategy="most_frequent")),
    ("onehot", OneHotEncoder(handle_unknown="ignore",
sparse_output=False))
])

preprocessor = ColumnTransformer([
    ("num", numeric_pipe, num_cols),
    ("cat", categorical_pipe, cat_cols)
])

Numeric columns used: ['no_of_dependents', 'income_annum',
'loan_amount', 'loan_term', 'cibil_score', 'residential_assets_value',
'commercial_assets_value', 'luxury_assets_value', 'bank_asset_value']
Categorical columns used: ['education', 'self_employed']

# 7. Model Training & Evaluation

models = {
    "Logistic Regression": LogisticRegression(max_iter=1000,
class_weight="balanced"),
    "Random Forest": RandomForestClassifier(n_estimators=200,
random_state=RND, class_weight="balanced"),
    "Gradient Boosting": GradientBoostingClassifier(random_state=RND)
}

roc_curves = {}

for name, model in models.items():
    pipe = Pipeline([("prep", preprocessor), ("model", model)])
    pipe.fit(X_train, y_train)

    preds = pipe.predict(X_test)
    proba = pipe.predict_proba(X_test)[:, 1] if hasattr(pipe,
"predict_proba") else None

```

```

print(f"\n{name} Results:")
print(classification_report(y_test, preds))

# Confusion matrix
sns.heatmap(confusion_matrix(y_test, preds), annot=True, fmt="d",
cmap="Blues")
plt.title(f"Confusion Matrix - {name}")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

if proba is not None:
    fpr, tpr, _ = roc_curve(y_test, proba)
    roc_auc = auc(fpr, tpr)
    roc_curves[name] = (fpr, tpr, roc_auc)

# ROC curves plot
plt.figure(figsize=(6,5))
for name, (fpr, tpr, roc_auc) in roc_curves.items():
    plt.plot(fpr, tpr, label=f"{name} (AUC={roc_auc:.2f})")

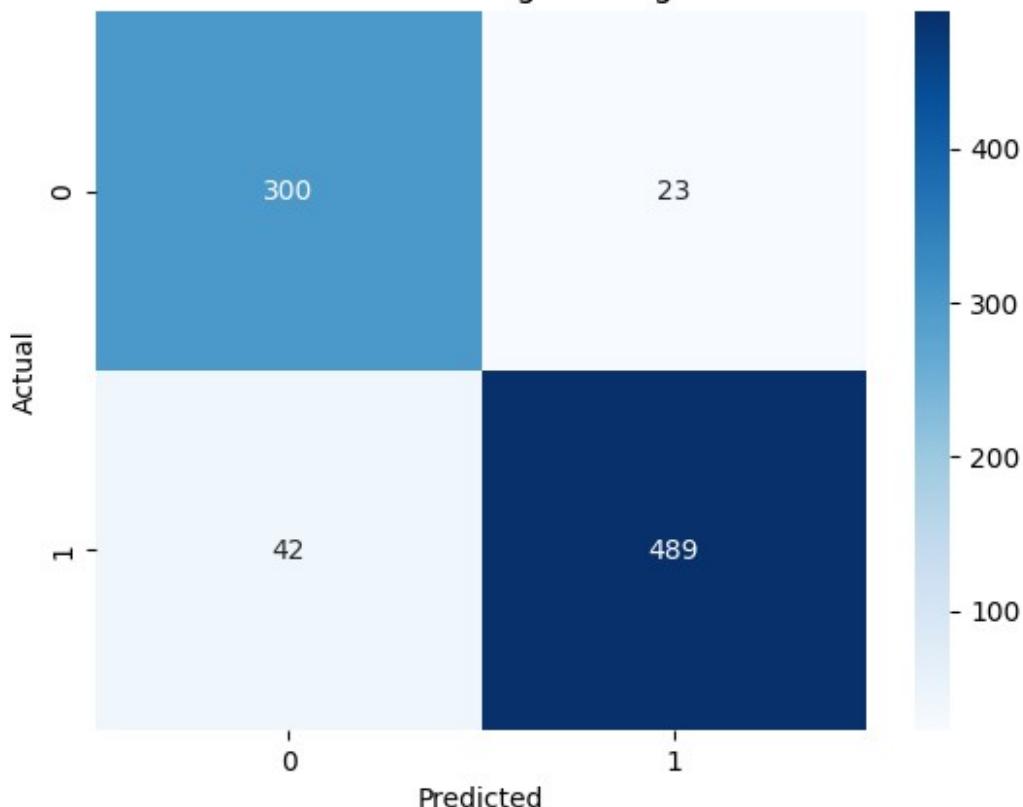
plt.plot([0,1], [0,1], "k--", label="Random guess")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curves")
plt.legend()
plt.tight_layout()
plt.show()

```

Logistic Regression Results:

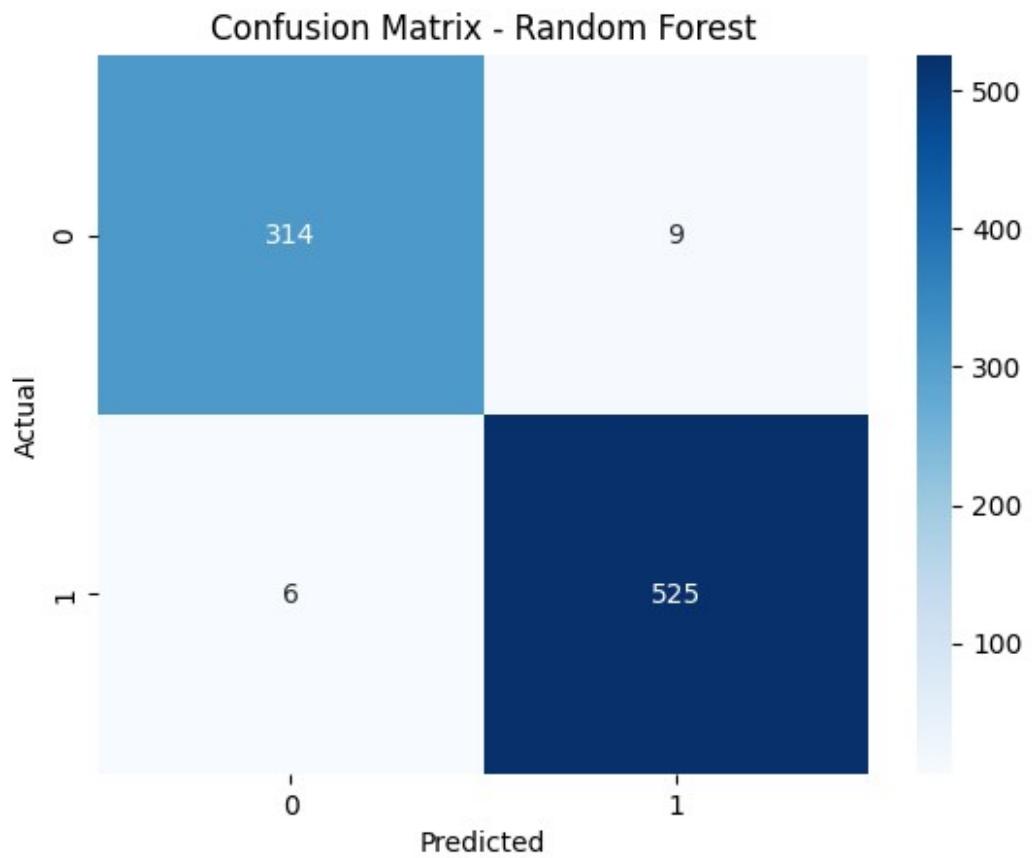
	precision	recall	f1-score	support
0	0.88	0.93	0.90	323
1	0.96	0.92	0.94	531
accuracy			0.92	854
macro avg	0.92	0.92	0.92	854
weighted avg	0.93	0.92	0.92	854

Confusion Matrix - Logistic Regression



Random Forest Results:

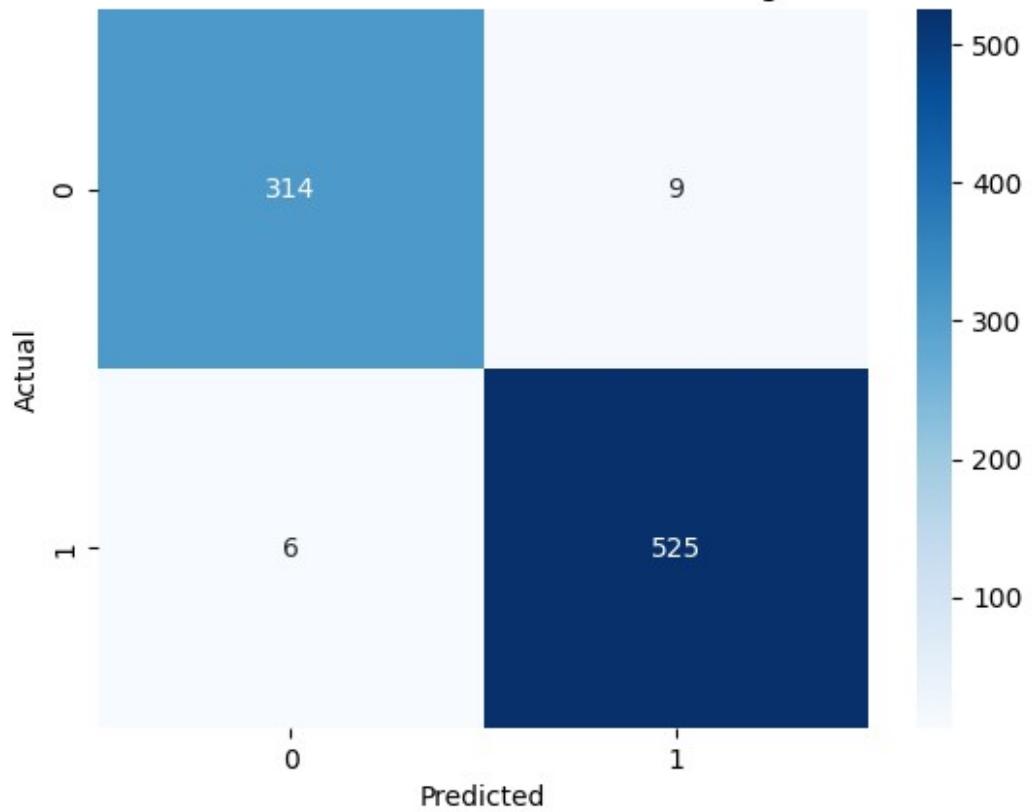
	precision	recall	f1-score	support
0	0.98	0.97	0.98	323
1	0.98	0.99	0.99	531
accuracy			0.98	854
macro avg	0.98	0.98	0.98	854
weighted avg	0.98	0.98	0.98	854

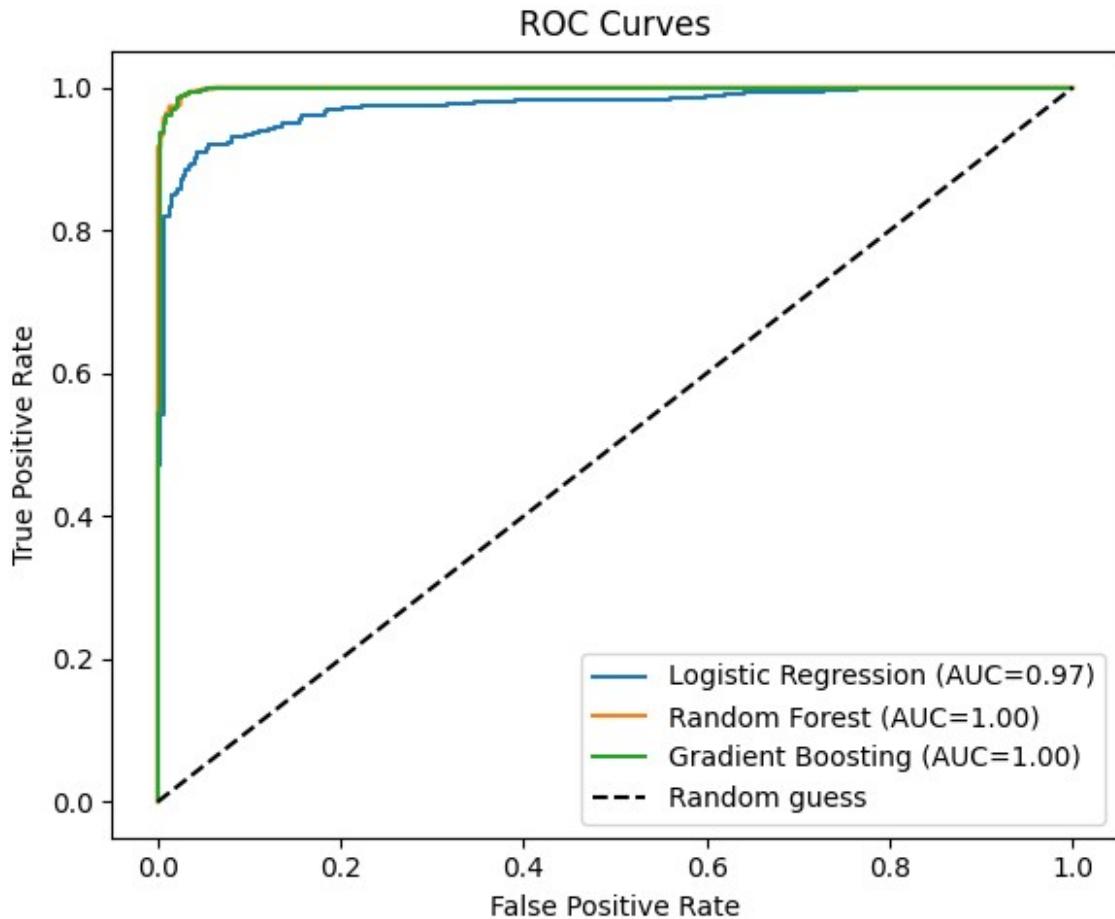


Gradient Boosting Results:

	precision	recall	f1-score	support
0	0.98	0.97	0.98	323
1	0.98	0.99	0.99	531
accuracy			0.98	854
macro avg	0.98	0.98	0.98	854
weighted avg	0.98	0.98	0.98	854

Confusion Matrix - Gradient Boosting





```

# =====
# 8. Feature Importance (Random Forest)

best_model = RandomForestClassifier(
    n_estimators=200, random_state=RND, class_weight="balanced"
)
pipe_best = Pipeline([("prep", preprocessor), ("model", best_model)])
pipe_best.fit(X_train, y_train)

# need fitted preprocessor to get OHE feature names
preprocessor.fit(X_train)

# numeric + one-hot encoded categorical feature names
num_feature_names = list(num_cols)
ohe = preprocessor.named_transformers_["cat"].named_steps["onehot"]
cat_feature_names = ohe.get_feature_names_out(cat_cols)
all_feature_names = num_feature_names + list(cat_feature_names)

importances = pipe_best.named_steps["model"].feature_importances_

```

```

fi = pd.DataFrame({"Feature": all_feature_names, "Importance": importances})
fi = fi.sort_values("Importance", ascending=False).head(15)

plt.figure(figsize=(8,6))
sns.barplot(x="Importance", y="Feature", data=fi)
plt.title("Top Feature Importances (Random Forest)")
plt.tight_layout()
plt.show()

```

